

Name: Priansh Madan

Batch: E4

Roll no.: 58

1. Perform logistic regression on the admission dataset

a) Import the Admission_Predict dataset, display it, and visualize various columns:

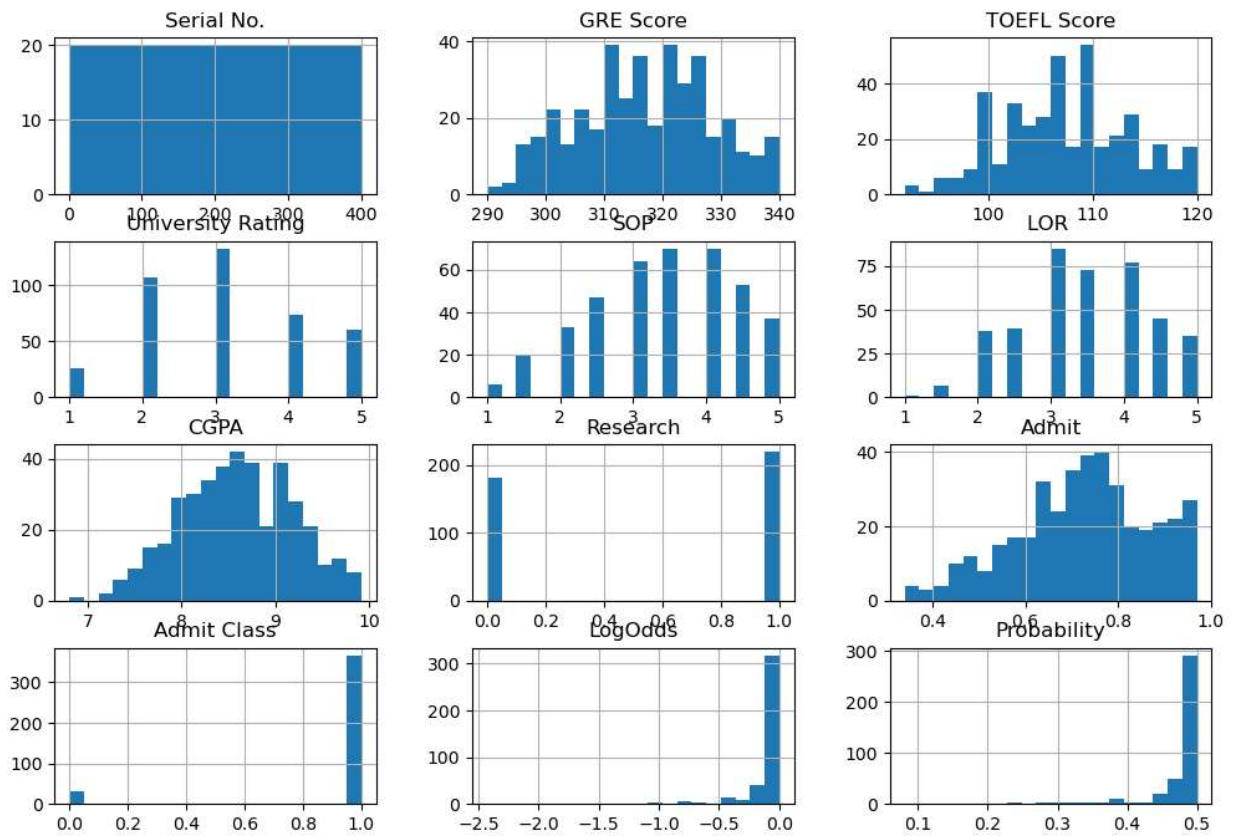
```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

data = pd.read_csv("Admission_Predict (1).csv")

print(data.head())
data.hist(bins=20, figsize=(12, 8))
plt.show()
```

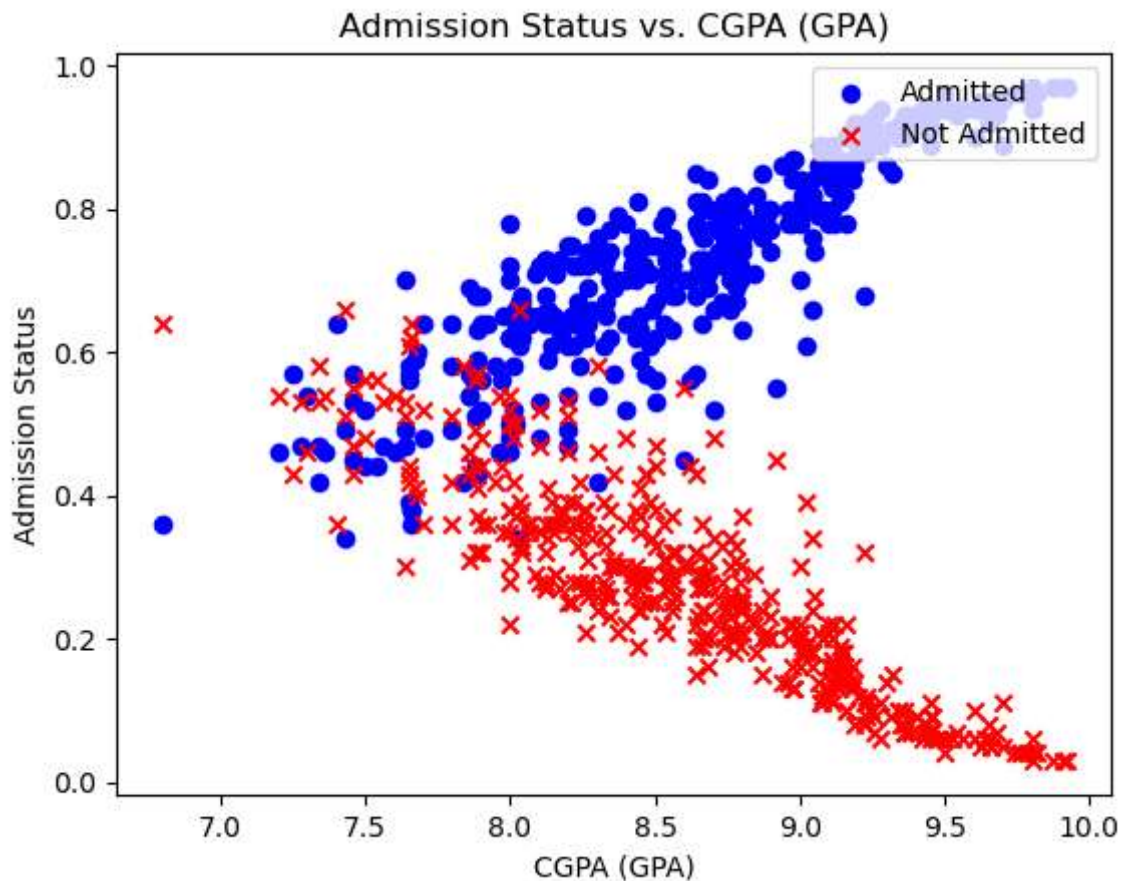
	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	\
0	1	337	118	4	4.5	4.5	9.65	
1	2	324	107	4	4.0	4.5	8.87	
2	3	316	104	3	3.0	3.5	8.00	
3	4	322	110	3	3.5	2.5	8.67	
4	5	314	103	2	2.0	3.0	8.21	

	Research	Admit	Admit Class	LogOdds	Probability
0	1	0.92	1	-0.000671	0.499832
1	1	0.76	1	-0.009353	0.497662
2	1	0.72	1	-0.164149	0.459055
3	1	0.80	1	-0.018318	0.495421
4	0	0.65	1	-0.084027	0.479006



b) Plot the dataset on GPA vs. admit score:

```
In [ ]: # b) Plot the dataset on GPA vs. admit score.
plt.scatter(data['CGPA'], data['Admit'], marker='o', c='b', label='Admitted')
plt.scatter(data['CGPA'], 1 - data['Admit'], marker='x', c='r', label='Not Admitted')
plt.xlabel('CGPA (GPA)')
plt.ylabel('Admission Status')
plt.legend(loc='upper right')
plt.title('Admission Status vs. CGPA (GPA)')
plt.show()
```



Find the slope and intercept of the line to fit:

```
In [ ]: data['Admit Class'] = (data['Admit'] >= 0.5).astype(int)
data.to_csv('Admission_Predict.csv', index=False)

X = data[['CGPA']].values # Reshape feature to (n_samples, n_features)
y = data['Admit Class'].values # Reshape target to (n_samples, )
X=X.reshape(-1,1)
model = LogisticRegression()
model.fit(X, y)

slope = model.coef_[0][0]
intercept = model.intercept_[0]
print(f"Slope: {slope}, Intercept: {intercept}")
```

Slope: 3.3834358054129954, Intercept: -25.343700528067235

Compute the log odds for each entry and merge the results with the data as a new column:

```
In [ ]: # Compute Log odds for each entry
log_odds = model.predict_log_proba(X)[:, 1]

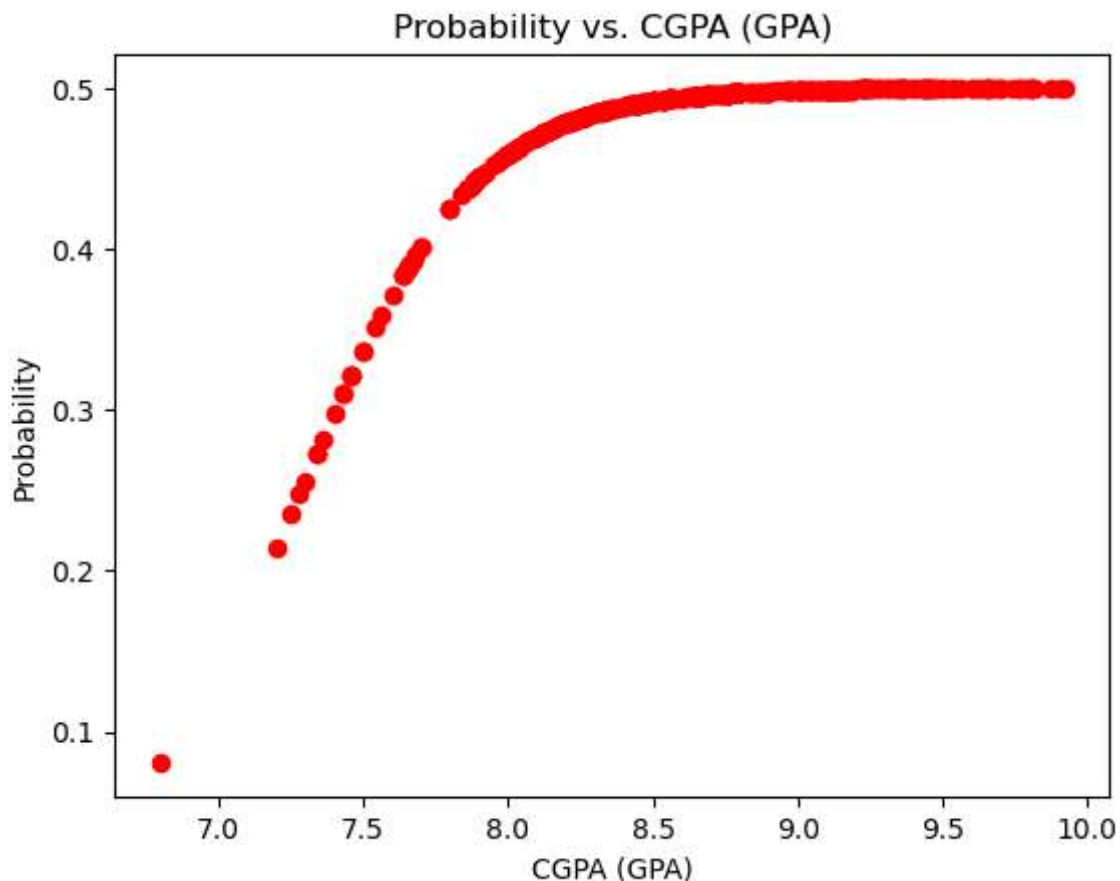
# Merge Log odds as a new column in the DataFrame
data['LogOdds'] = log_odds
```

e) Using the log odds compute the probability for each entry.

```
In [ ]: data['Probability'] = 1 / (1 + np.exp(-log_odds))
data.to_csv('Admission_Predict (1).csv', index=False)
```


f) Plot the probabilities vs gpa graph.

```
In [ ]: plt.scatter(data['CGPA'], data['Probability'], marker='o', c='r', label='Probability')
plt.xlabel('CGPA (GPA)')
plt.ylabel('Probability')
plt.title('Probability vs. CGPA (GPA)')
plt.show()
```



g) Show the decision boundary of the regression model.

```
In [ ]: X_values = np.linspace(data['CGPA'].min(), data['CGPA'].max(), 100)
decision_boundary = -intercept / slope # Decision boundary where log odds = 0
plt.scatter(data['CGPA'], data['Admit'], marker='o', c='b', label='Admitted')
plt.scatter(data['CGPA'], 1 - data['Admit'], marker='x', c='g', label='Not Admitted')
plt.axvline(decision_boundary, color='k', linestyle='--', label='Decision Boundary')
plt.xlabel('CGPA (GPA)')
plt.ylabel('Admission Status')
plt.legend(loc='upper right')
plt.title('Admission Status vs. CGPA (GPA) with Decision Boundary')
plt.show()
```



h) Show the accuracy of the regressor model.

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy of the Logistic Regression Model: {accuracy * 100:.2f}%")
```

Accuracy of the Logistic Regression Model: 89.17%

In []:

2. Perform logistic regression on the credit card dataset

Load the dataset, visualize it, show the data headers

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, precision_score, recall_score

# Load the dataset
df = pd.read_csv('creditcard.csv')
```

```
# Show the first few rows of the dataset
print(df.head())


# Visualize the class distribution (fraudulent vs. non-fraudulent)
plt.figure(figsize=(8, 6))
df['Class'].value_counts().plot(kind='bar', color=['skyblue', 'salmon'])
plt.title('Class Distribution (0: Non-Fraudulent, 1: Fraudulent)')
plt.xlabel('Class')
plt.ylabel('Count')
plt.show()
```

	Time	V1	V2	V3	V4	V5	V6	V7	\
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	

	V8	V9	...	V21	V22	V23	V24	V25	\
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170	
2	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642	
3	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	
4	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010	

	V26	V27	V28	Amount	Class
0	-0.189115	0.133558	-0.021053	149.62	0
1	0.125895	-0.008983	0.014724	2.69	0
2	-0.139097	-0.055353	-0.059752	378.66	0
3	-0.221929	0.062723	0.061458	123.50	0
4	0.502292	0.219422	0.215153	69.99	0

[5 rows x 31 columns]



c) Preprocess the dataset if required

```
In [ ]: # i. Check for duplicate data and remove it:
df = df.drop_duplicates()
print("Number of duplicate rows:", df.duplicated().sum())
```

Number of duplicate rows: 0

```
In [ ]: # Remove unnecessary columns like 'Time':
df = df.drop(columns=['Time'])
```

```
In [ ]: # Separate the dataset into feature and target columns:
X = df.drop(columns=['Class']) # Features
y = df['Class'] # Target
```

```
In [ ]: # Scale the dataset using standard scaling:
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
In [ ]: # Partition the dataset into training and testing sets (80%-20%):

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

# Check the shapes of the resulting datasets
```



```
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
```

X_train shape: (226980, 29)

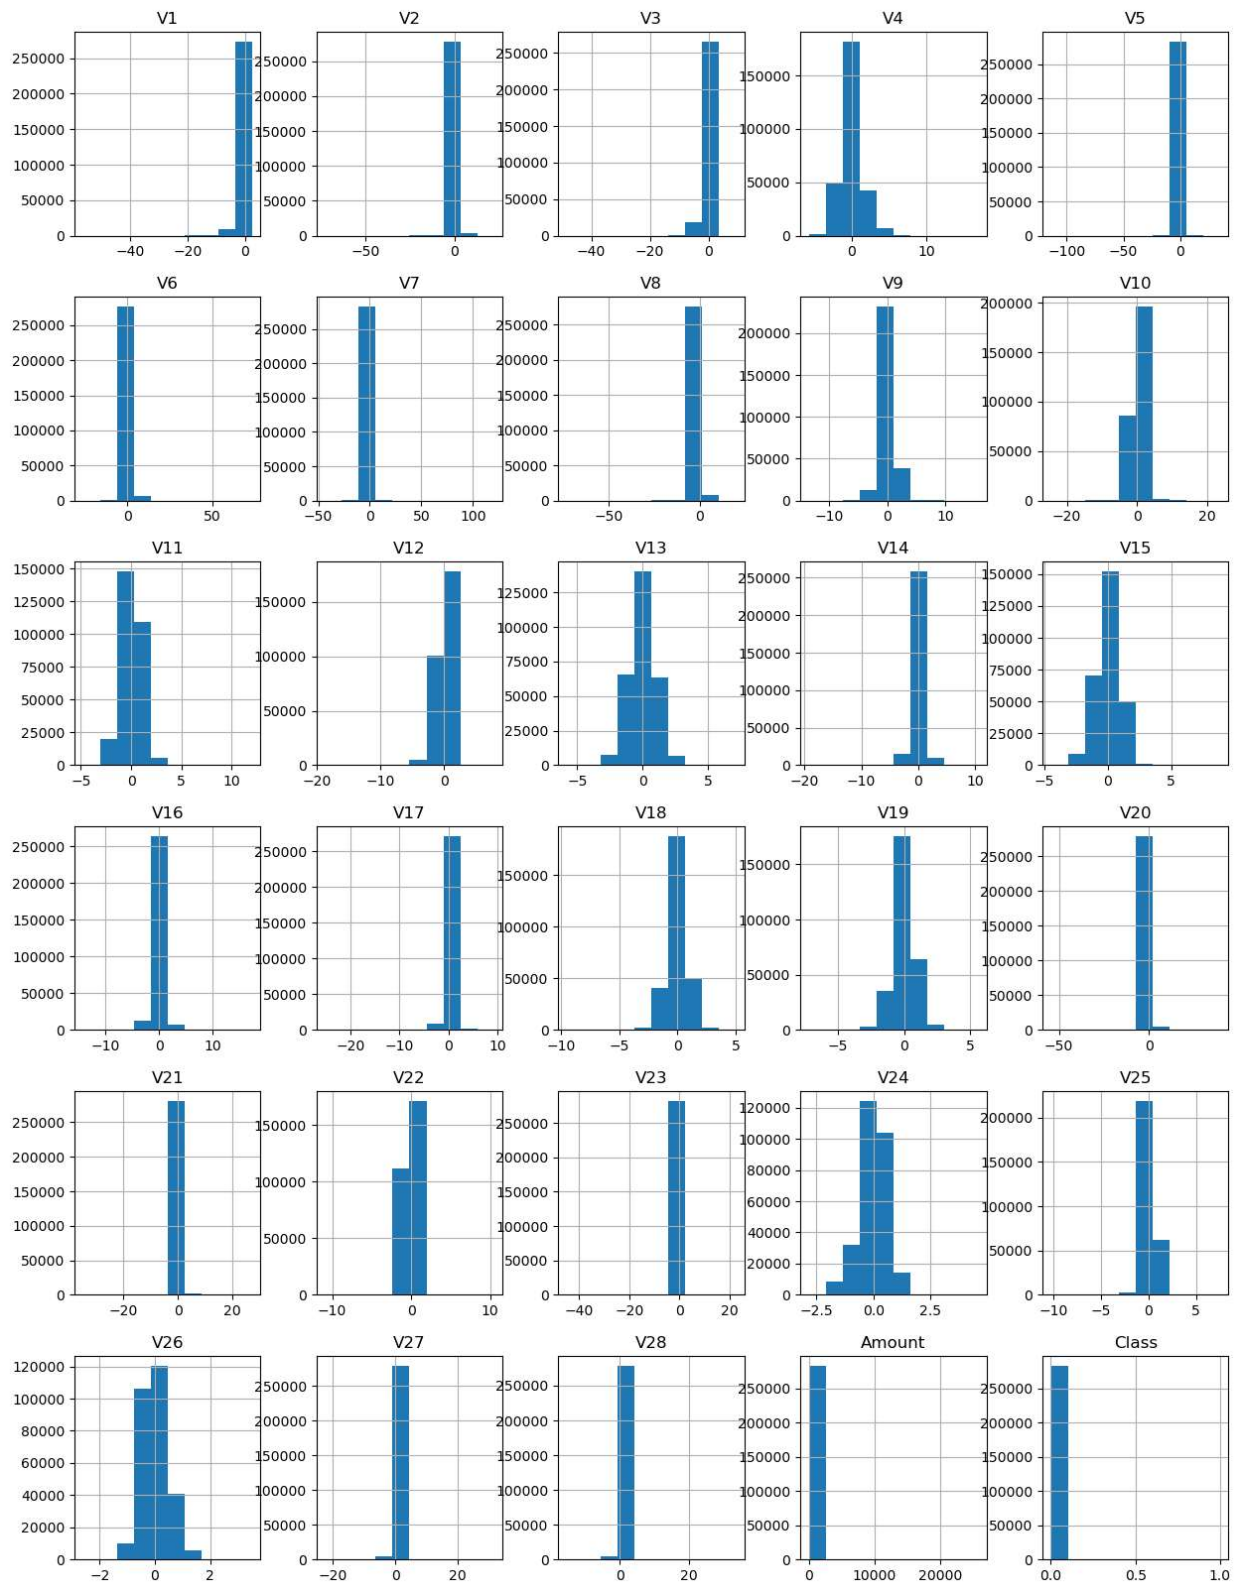
X_test shape: (56746, 29)


y_train shape: (226980,)

y_test shape: (56746,)

Step d: Plot histograms/heatmaps

```
In [ ]: df.hist(figsize=(15, 20))
plt.show()
corr = df.corr()
plt.figure(figsize=(12, 10))
sns.heatmap(corr, annot=True, cmap='coolwarm')
plt.show()
```



Step e: Train the model using logistic regression

```
In [ ]: model = LogisticRegression()
        model.fit(X_train, y_train)
```

```
Out[ ]: LogisticRegression()
```

Step f: Obtain the training accuracy

```
In [ ]: train_accuracy = model.score(X_train, y_train)
        print(f"Training Accuracy: {train_accuracy}")
```

Training Accuracy: 0.9991937615648956

Step g: Test the model and obtain the testing accuracy

```
In [ ]: test_accuracy = model.score(X_test, y_test)
        print(f"Testing Accuracy: {test_accuracy}")
```

Testing Accuracy: 0.9991717477883904

Step h: Generate confusion matrix, precision and recall

```
In [ ]: y_pred = model.predict(X_test)
        confusion = confusion_matrix(y_test, y_pred)
        precision = precision_score(y_test, y_pred)
        recall = recall_score(y_test, y_pred)
```

```
In [ ]: print(f"Confusion Matrix:\n{confusion}")
        print(f"Precision: {precision}")
        print(f"Recall: {recall}")
```

```
Confusion Matrix:
[[56650   6]
 [  41   49]]
Precision: 0.8909090909090909
Recall: 0.5444444444444444
```

```
In [ ]:
```